

Performance

A20379025 Hsueh Yi Chen
A20385942 Yitong Huang

CPU Benchmark:

- The program finds GFLOPS and GIOPS per second
- The start time and end time is calculated
- The total time is difference between start time and end time
- The performance is measured using the formula

$$\text{OPS} = \text{number of instructions} / \text{Total time}$$

System Configuration:

1. for without AVX program and 600 sample program

Chameleon Cloud(openstack KVM)

Processor: Intel Xeon E312xx (Sandy Bridge) @ 3.09 GHz

1 processor, 32 cores

OS:CentOS7

RAM: 4GB

Disk: 40GB

2. for without AVX program

Chameleon Cloud(baremetal)

Processor: Intel(R) Xeon(R) CPU E5-2670 v3 @ 2.30GHz


| Image Overview | |
|------------------|---|
| Information | |
| Name | PA1_ys |
| Description | https://www.chameleoncloud.org/appliances/1/0ec0a4a9-d76d-42c9-a4e6-7983919b2328 |
| ID | 0ec0a4a9-d76d-42c9-a4e6-7983919b2328 |
| Owner | CH-819402 |
| Status | Active |
| Public | No |
| Protected | No |
| Checksum | 14da33f3009cc620ea4f3f51dd2f7eb9 |
| Created | Oct. 9, 2017, 1:35 a.m. |
| Updated | Oct. 9, 2017, 1:37 a.m. |
| Specs | |
| Size | 1.8 GB |
| Container Format | BARE |
| Disk Format | QCOW2 |
| Min Disk | 80GB |

Experiment 1~8:

Two programs are run 1 time to calculate GIOPS and GFLOPS with AVX instructions and without AVX instructions, respectively. The result is recorded in the table 1 and table 2 shown below:

| Operations/sec | GIOPS | GFLOPS |
|----------------|----------|----------|
| 1 thread | 1.012469 | 1.012507 |
| 2 threads | 2.023625 | 2.023738 |
| 4 threads | 3.763475 | 3.765623 |
| 8 threads | 7.250366 | 7.2515 |


Table 1. Performance with AVX instructions



| Operations/sec | GIOPS | GFLOPS |
|----------------|----------|----------|
| 1 thread | 0.846123 | 0.863293 |
| 2 threads | 1.634691 | 1.680904 |
| 4 threads | 1.598174 | 1.606577 |

| | | |
|-----------|----------|----------|
| 8 threads | 1.624932 | 1.643697 |
|-----------|----------|----------|

Table 2. Performance without AVX instructions



Observation:

- After using AVX instructions, the efficiency of GFLOPS and GIOPS is 2 or 3 times better than the result without AVX instructions.
- Based on the result without AVX instructions, GIOPS and GFLOPS are always growing up when having 2 threads.
- In case of the result with AVX instructions, there is a constant linear growth as the number of threads increases. On the other hand, the result without AVX instructions continues to rise until hitting a plateau after getting more than 2 threads.

Experiment 9,10:

A program was run for 600 seconds and recorded 600 samples for IOPS and FLOPS. Each Dataset is plotted in the line chart 1 and chart 2. The charts are shown below:

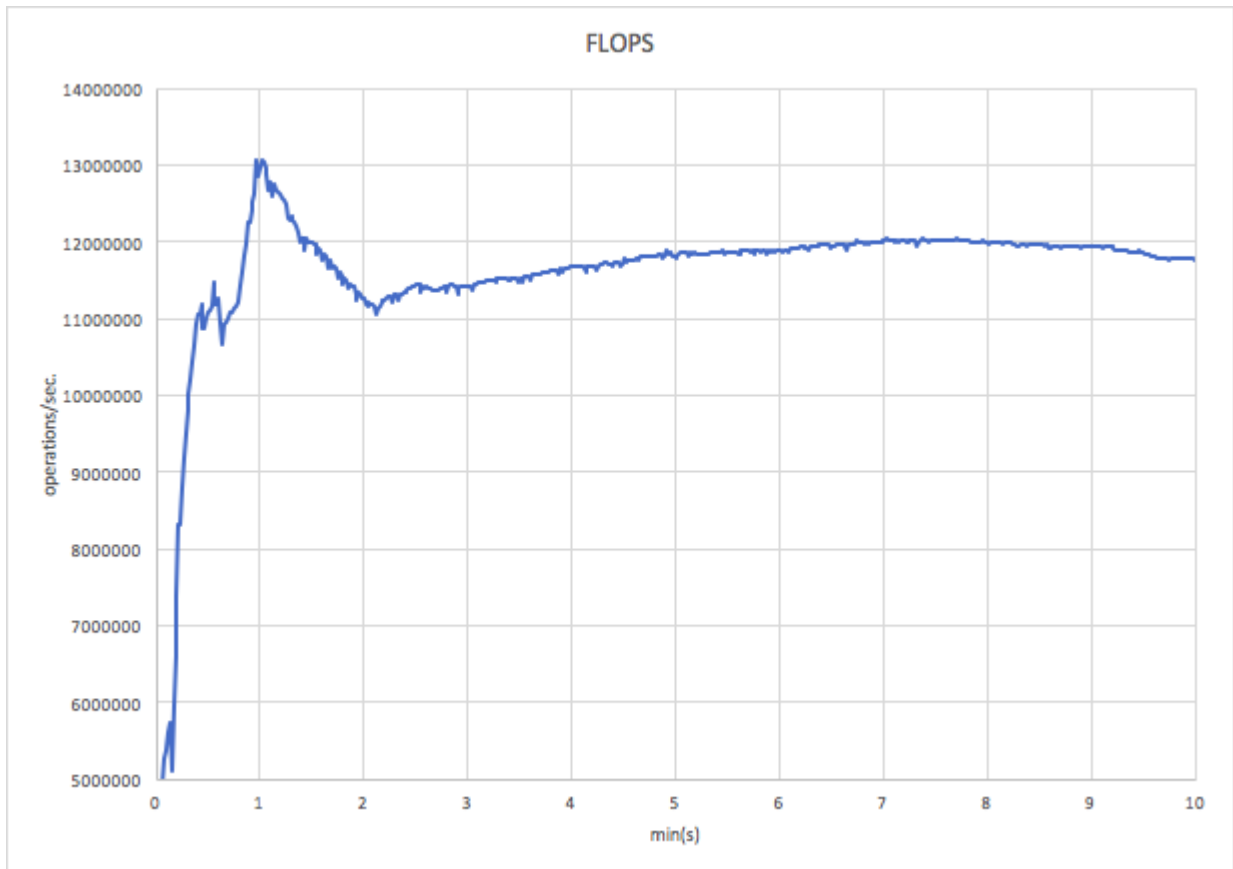


Chart 1.600 samples for FLOPS

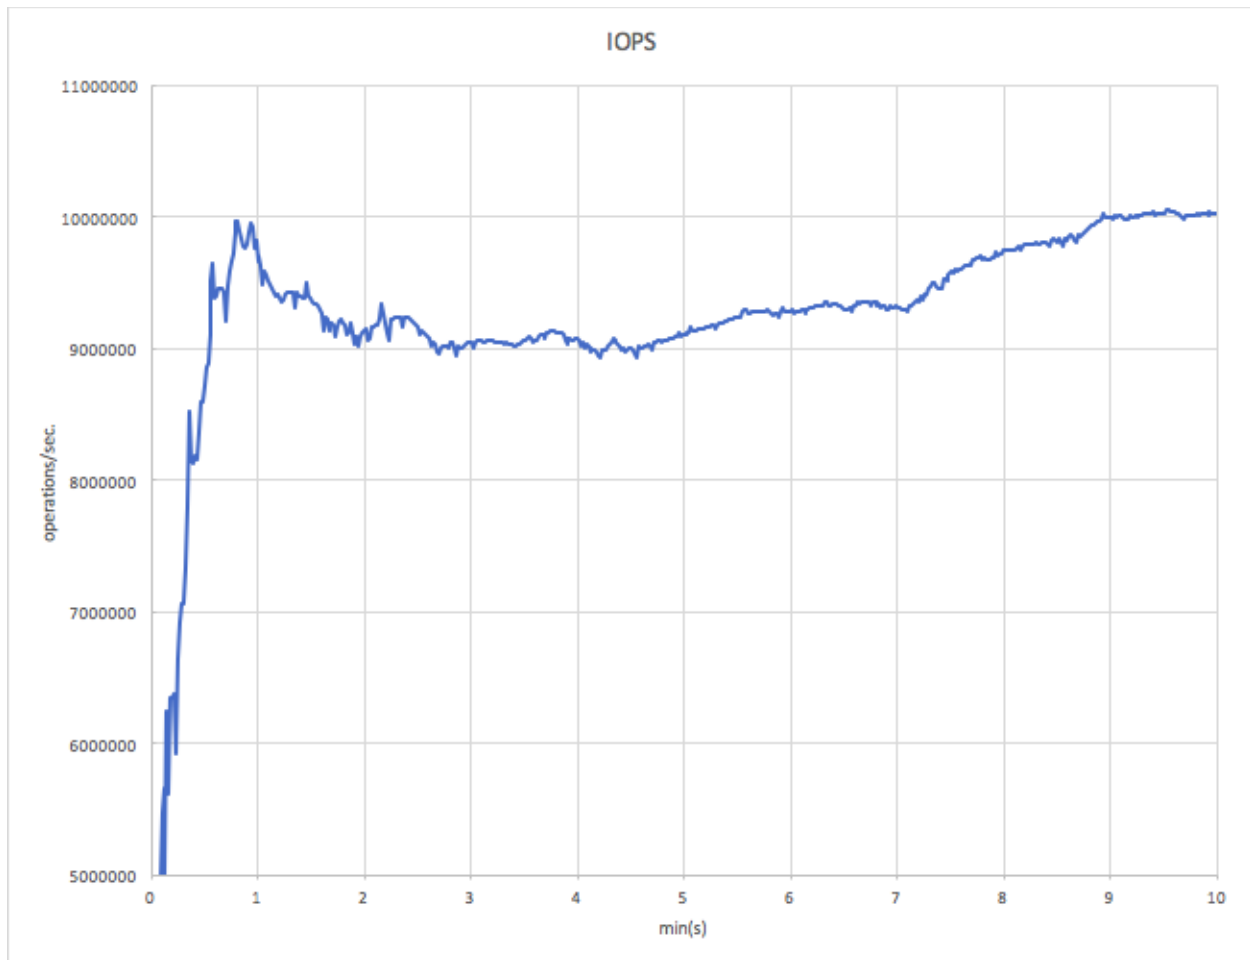


Chart 2.600 samples for IOPS

Observation:

- In the case of the result for IOPS and FLOPS, each of them has an apparent shooting up in the first minute. But in the second minute, each of them declines a little simultaneously. Then both of them upturn gradually in final eight minutes.
- The difference between IOPS and FLOPS is each sample is 10 million times larger than each sample in IOPS.

Experiment 11:

Linpack was run and the following results was shown below:

```

[cc@zxc linpack]$ ./xlinpack_xeon64
Input data or print help ? Type [data]/help :

Number of equations to solve (problem size): 2000
Leading dimension of array: 20000
Number of trials to run: ^C
[cc@zxc linpack]$ ./xlinpack_xeon64
Input data or print help ? Type [data]/help :
[data]/help
Number of equations to solve (problem size): 2000
Leading dimension of array: 2000
Number of trials to run: 4
Data alignment value (in Kbytes): 4
Current date/time: Sun Oct  8 06:02:37 2017

CPU frequency:    3.090 GHz
Number of CPUs: 2
Number of cores: 2
Number of threads: 2

Parameters are set to:

Number of tests: 1
Number of equations to solve (problem size) : 2000
Leading dimension of array                  : 2000
Number of trials to run                     : 4
Data alignment value (in Kbytes)           : 4

Maximum memory requested that can be used=32044096, at the size=2000

===== Timing linear equation system solver =====

Size   LDA   Align. Time(s)   GFlops   Residual   Residual(norm) Check
2000   2000   4      0.085    62.5099   3.329198e-12 2.895994e-02 pass
2000   2000   4      0.084    63.6462   3.329198e-12 2.895994e-02 pass
2000   2000   4      0.084    63.9332   3.329198e-12 2.895994e-02 pass
2000   2000   4      0.084    63.2805   3.329198e-12 2.895994e-02 pass

Performance Summary (GFlops)

Size   LDA   Align. Average Maximal
2000   2000   4      63.3425 63.9332

Residual checks PASSED

End of tests

```

Theoretical Peak Performance:

The CentOS Linux instance has a processor: Intel Xeon E312xx (Sandy Bridge) @3.09GHz
The closest to it is Intel® Xeon® Processor E3-1225 and its performance is :

| Processor number | Frequency Type | Clock | GFLOP |
|------------------|----------------|---------|-------|
| E3-1225 | Base | 3.10GHz | 99.2 |
| | Max Turbo | 3.40GHz | 109 |

Conclusion:

- Compared to theoretical Peak performance, the algorithm used in my own program is not as efficient as theoretical Peak performance. It wastes much time for lots of processing like assigning variables or passing pointer.
- In the case of Linpack performance, values of GFLOPS are much larger than values from my own program. Although I uses AVX instructions to improve the efficiency, but the difference of GFLOPS between Linpack and my own program is still large.

Memory Benchmark:

- The program measures the latency and throughput of memory speed by allocating 1.28GB memory
- There are three operations, sequential read+write, sequential write and random write
- Latency is measured in microseconds/8B
- Throughput is measured in MB/s

System Configuration:

Chameleon Cloud(openstack KVM)

Processor: Intel Xeon E312xx (Sandy Bridge) @ 3.09 GHz

OS:CentOS7

RAM: 16GB


Disk: 160GB

(instance size:xlarge, since it will get out of memory when running in medium size)

Experiment 1~48:


- Sequential read+write:
Throughput:

| Block size | 1 thread | 2 threads | 4 threads | 8 threads |
|------------|-------------|-------------|-------------|-------------|
| 8B | 1753.866728 | 3137.485599 | 5376.840938 | 9487.806686 |
| 8KB | 4992.530706 | 9448.098201 | 17447.48715 | 32504.63445 |
| 8MB | 5529.969844 | 9299.89247 | 16986.26501 | 29970.96563 |
| 80MB | 5549.77454 | 9638.917128 | 17392.72223 | 32482.36309 |



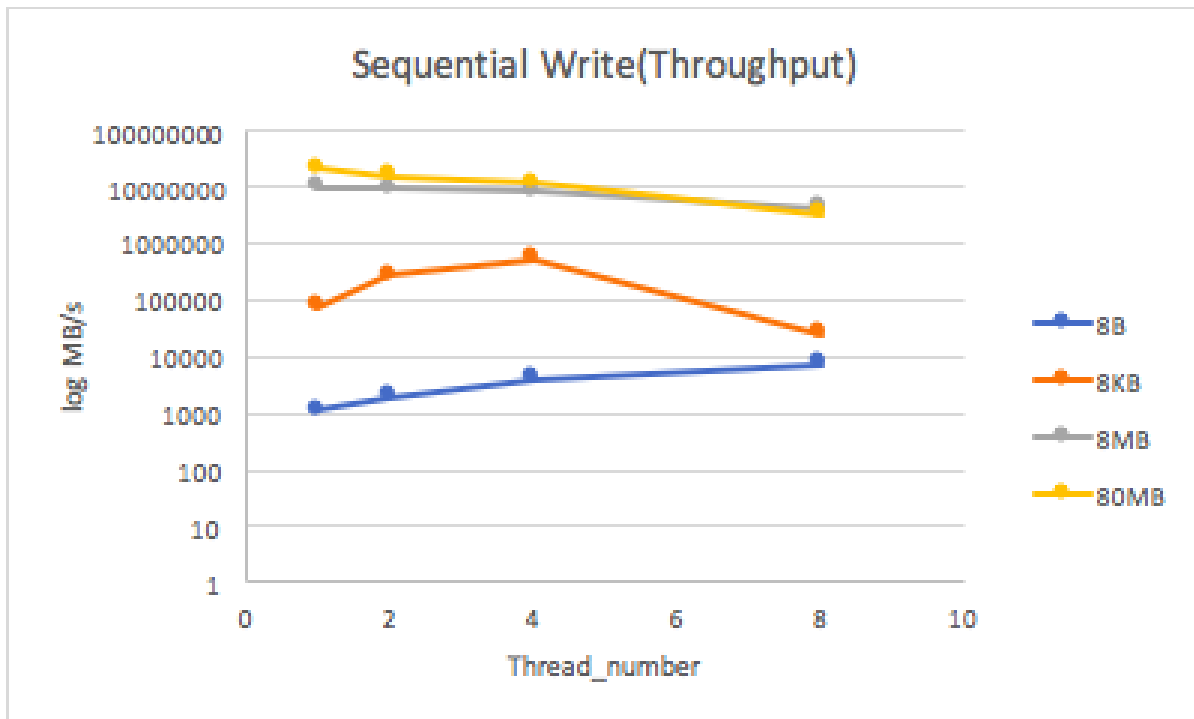
Latency:

| Block size | 1 thread | 2 threads | 4 threads | 8 threads |
|------------|-------------|-------------|-------------|-------------|
| 8B | 0.00456135 | 0.002549812 | 0.001487862 | 0.000843187 |
| 8KB | 0.001602394 | 0.000846731 | 0.000458519 | 0.000246119 |
| 8MB | 0.001446662 | 0.000860225 | 0.000470969 | 0.000266925 |
| 80MB | 0.0014415 | 0.000829969 | 0.000459963 | 0.000246287 |



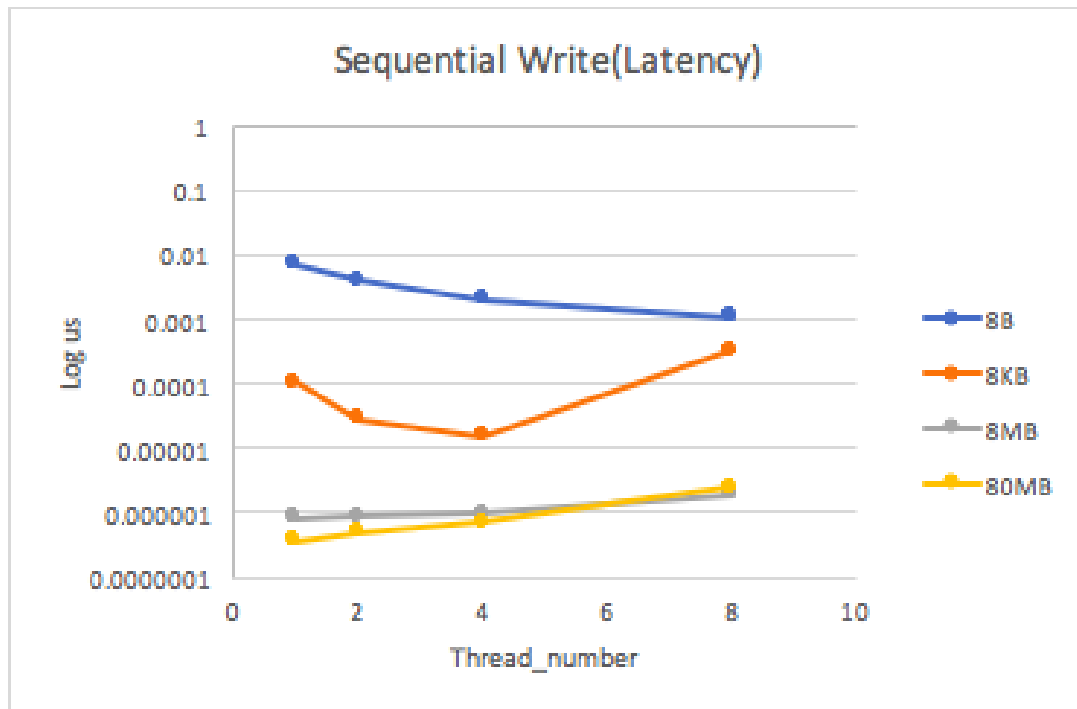
- Sequential Read:
Throughput:

| Block size | 1 thread | 2 threads | 4 threads | 8 threads |
|------------|-------------|-------------|-------------|-------------|
| 8B | 1140.350408 | 1969.912662 | 3947.620017 | 7419.257611 |
| 8KB | 77122.37151 | 276996.3211 | 521810.0285 | 24931.34142 |
| 8MB | 9624060.15 | 9343065.693 | 8258064.516 | 4155844.156 |
| 80MB | 21694915.25 | 16202531.65 | 11531531.53 | 3224181.36 |



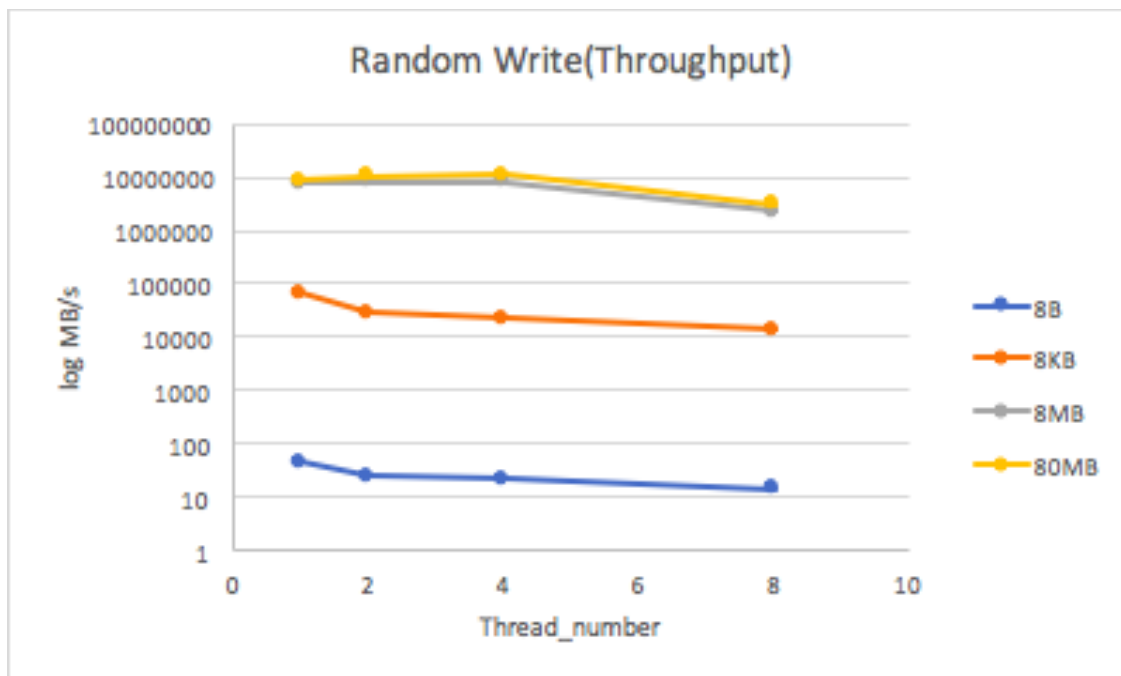
Latency:

| Block size | 1 thread | 2 threads | 4 threads | 8 threads |
|------------|-------------|-------------|-------------|-------------|
| 8B | 0.007015388 | 0.004061094 | 0.002026537 | 0.001078275 |
| 8KB | 0.000103731 | 0.000028881 | 0.000015331 | 0.000320881 |
| 8MB | 0.000000831 | 0.000000856 | 0.000000969 | 0.000001925 |
| 80MB | 0.000000369 | 0.000000494 | 0.000000694 | 0.000002481 |




- Random Write:
Throughput:

| Block size | 1 thread | 2 threads | 4 threads | 8 threads |
|------------|-------------|-------------|-------------|-------------|
| 8B | 46.189383 | 25.173689 | 22.304231 | 14.148185 |
| 8KB | 66273.16972 | 27964.69457 | 21829.96504 | 13555.15784 |
| 8MB | 7619047.619 | 8476821 | 8152866.242 | 2335766.423 |
| 80MB | 9014084.507 | 10322580.65 | 11428571.43 | 3033175 |



Latency:

| Block size | 1 thread | 2 threads | 4 threads | 8 threads |
|------------|-------------|-------------|-------------|-------------|
| 8B | 0.173199975 | 0.317792119 | 0.358676344 | 0.565443563 |
| 8KB | 0.000120713 | 0.000286075 | 0.000366469 | 0.000590181 |
| 8MB | 0.00000105 | 0.000000944 | 0.000000981 | 0.000003425 |
| 80MB | 0.000000887 | 0.000000775 | 0.0000007 | 0.000002638 |



Observation:

- As the number of thread grows up, the speed of transfer also increases until hitting a plateau when having 4 threads. Some operations like random write and sequential write even get worse efficiency when having 8 threads.
- Sequential write is always faster than random write
- Increasing the block size can't always improve the speed of transfer. In most of experiments, 8MB and 80MB almost get the same performance.
- Small block size like 8B always get the highest latency.
- Based on the result of each operations, 4 threads of concurrency can get the best performance.

Experiment 49:

Stream was run to estimate the memory benchmarking

```
[cc@zxc ~]$ gcc stream.c
[cc@zxc ~]$ ./a.out

-----
STREAM version $Revision: 5.10 $
-----

This system uses 8 bytes per array element.
-----

Array size = 10000000 (elements), Offset = 0 (elements)
Memory per array = 76.3 MiB (= 0.1 GiB).
Total memory required = 228.9 MiB (= 0.2 GiB).
Each kernel will be executed 10 times.
The *best* time for each kernel (excluding the first iteration)
will be used to compute the reported bandwidth.
-----

Your clock granularity/precision appears to be 1 microseconds.
Each test below will take on the order of 36774 microseconds.
(= 36774 clock ticks)
Increase the size of the arrays if this shows that
you are not getting at least 20 clock ticks per test.
-----

WARNING -- The above is only a rough guideline.
For best results, please be sure you know the
precision of your system timer.
-----

Function      Best Rate MB/s  Avg time     Min time     Max time
Copy:         5900.8   0.028029    0.027115    0.029005
Scale:        5727.4   0.028919    0.027936    0.030539
Add:          8378.1   0.029911    0.028646    0.033670
Triad:        7964.7   0.032292    0.030133    0.036875
-----

Solution Validates: avg error less than 1.000000e-13 on all three arrays
-----
```

Theoretical Peak Performance:

The CentOs Linux instance has a processor: Intel Xeon E312xx (Sandy Bridge) 3.09 GHz
RAM:16GB

The closet to its is Intel Xeon E312xx (Sandy Bridge) @ 1.80 GHz RAM:32GB and its
performance is :

| Operation type | Speed |
|-----------------|----------|
| Sequential Read | 3.52GB/s |
| Sequential | 2.94GB/s |

| | |
|-------------|----------|
| Write | |
| Stdlib Copy | 1.96GB/s |

Conclusion:

- The formulated values of each operation mentioned above are much larger than result in Stream benchmark and theoretical performance.
- Compared to theoretical memory bandwidth and Stream benchmark, the algorithm used in my own program missed some part of memory when allocating memory. Therefore, the actual values would be less than formulated values.

Disk Benchmark

- The program finds throughput in MB per second and latency in ms
- The 10 GB binary file is allocated first
- The start time and end time is calculated only for actual read+rewrite, sequential read and random read
- The total time is difference between start and end time
- The actual operation data size is adjusted during experiment to make sure the run time is at least 10 seconds but no so long

- The throughput performance is measured using formula

$$\text{Throughput} = \text{total operation data size} / \text{total time}$$

- The latency performance is measured using formula

$$\text{Latency} = \text{total time} / \text{total operation data size} \times 8 \text{ (since requirement asked to use 8B block size to measure latency)}$$

System Configuration

Chameleon Cloud(openstack KVM)

Processor: Intel Xeon E312xx (Sandy Bridge) @ 3.09 GHz

1 processor, 32 cores

OS:CentOS7

RAM: 4GB

Disk: 40GB

Experiment 1~16


- Read Sequential program is run for testing sequential read performance within the 10 GB binary data file. The sequential read data size in below table means the data size operated by each experiment. This data size is adjusted during testing to make sure the run time is 10s. But the sequential read data size was fixed for each block size to make sure strong scaling.


Result shows in below table and graphs.

| block_size | thread number | throughput (MBps) | latency(ms) | run time(s) | Sequential Read | file datasize |
|------------|---------------|-------------------|--------------|-------------|-----------------|---------------|
|------------|---------------|-------------------|--------------|-------------|-----------------|---------------|

| | | | | | datasize(B) | |
|------|---|-------------|----------|-----------|-------------|------|
| 8B | 1 | 23.347399 | 0.000343 | 42.831324 | 1.00E+09 | 10GB |
| | 2 | 33.107913 | 0.000242 | 30.20426 | 1.00E+09 | 10GB |
| | 4 | 35.227066 | 0.000227 | 28.387263 | 1.00E+09 | 10GB |
| | 8 | 36.245878 | 0.000221 | 27.589344 | 1.00E+09 | 10GB |
| 8KB | 1 | 580.293798 | 0.000014 | 51.69795 | 3.00E+10 | 10GB |
| | 2 | 390.117129 | 0.000021 | 76.899981 | 3.00E+10 | 10GB |
| | 4 | 1145.140031 | 0.000007 | 26.197669 | 3.00E+10 | 10GB |
| | 8 | 1571.941037 | 0.000005 | 19.084685 | 3.00E+10 | 10GB |
| 8MB | 1 | 1675.124039 | 0.000005 | 29.848536 | 5.00E+10 | 10GB |
| | 2 | 1536.660188 | 0.000005 | 32.538098 | 5.00E+10 | 10GB |
| | 4 | 1762.40535 | 0.000005 | 28.370318 | 5.00E+10 | 10GB |
| | 8 | 1936.964278 | 0.000004 | 25.813589 | 5.00E+10 | 10GB |
| 80MB | 1 | 1137.94976 | 0.000007 | 43.938671 | 5.00E+10 | 10GB |
| | 2 | 806.951948 | 0.00001 | 61.961558 | 5.00E+10 | 10GB |
| | 4 | 1034.113558 | 0.000008 | 48.350589 | 5.00E+10 | 10GB |
| | 8 | 1774.294539 | 0.000005 | 28.180214 | 5.00E+10 | 10GB |

Table Performance of sequential read





Experiment 17~32



- Read Random program is run for testing Random read performance within the 10 GB binary data file.
- The randomization was implemented by rand() function. We seek a random block sized pointer before every reading. And then read the randomly seeking block.
- Other notifications are the same with experiment 1~16.

Result shows in below table and graphs.

| block_size | thread number | throughput (MBps) | latency(ms) | run time(s) | Random Read datasize(B) | file datasize |
|------------|---------------|-------------------|-------------|-------------|-------------------------|---------------|
| 8B | 1 | 12.625589 | 0.000634 | 79.204226 | 1.00E+09 | 10GB |
| | 2 | 18.104972 | 0.000442 | 55.233447 | 1.00E+09 | 10GB |
| | 4 | 15.836879 | 0.000505 | 63.143756 | 1.00E+09 | 10GB |
| | 8 | 18.266729 | 0.000438 | 54.744339 | 1.00E+09 | 10GB |
| 8KB | 1 | 4622.064161 | 0.000002 | 17.308284 | 8.00E+10 | 10GB |
| | 2 | 6836.899511 | 0.000001 | 11.70121 | 8.00E+10 | 10GB |
| | 4 | 7386.729248 | 0.000001 | 10.830233 | 8.00E+10 | 10GB |
| | 8 | 7576.294088 | 0.000001 | 10.559252 | 8.00E+10 | 10GB |
| 8MB | 1 | 6027.785294 | 0.000001 | 33.179682 | 2.00E+11 | 10GB |
| | 2 | 10124.20053 | 0.000001 | 19.754646 | 2.00E+11 | 10GB |
| | 4 | 9708.827844 | 0.000001 | 20.599809 | 2.00E+11 | 10GB |
| | 8 | 10318.74739 | 0.000001 | 18.715544 | 2.00E+11 | 10GB |
| 80MB | 1 | 3779.746878 | 0.000002 | 26.456798 | 1.00E+11 | 10GB |
| | 2 | 7090.083106 | 0.000001 | 14.104207 | 1.00E+11 | 10GB |
| | 4 | 6581.964264 | 0.000001 | 15.193033 | 1.00E+11 | 10GB |

| | | | | | | |
|--|---|-------------|----------|-----------|----------|------|
| | 8 | 7437.105123 | 0.000001 | 13.446092 | 1.00E+11 | 10GB |
|--|---|-------------|----------|-----------|----------|------|

Table Performance of random read



Experiment 33~48



- Read+write program is run for testing read and rewrite performance within the 10 GB binary data file.

- The program was implemented as following, first sequentially read the whole file; then sequentially re-write the file for a fixed data size.
- Other notifications are the same with experiment 1~16.

Result shows in below table and graphs.

| block_size | thread number | throughput (MBps) | latency(ms) | run time(s) | Read Datasize(B) | Write Datasize(B) | file Datasize |
|------------|---------------|-------------------|-------------|-------------|------------------|-------------------|---------------|
| 8B | 1 | 23.691809 | 0.000338 | 44.319115 | 1.00E+09 | 5.00E+08 | 10GB |
| | 2 | 30.948178 | 0.000258 | 33.927683 | 1.00E+09 | 5.00E+08 | 10GB |
| | 4 | 26.165928 | 0.000306 | 40.128521 | 1.00E+09 | 5.00E+08 | 10GB |
| | 8 | 34.01607 | 0.000235 | 30.867764 | 1.00E+09 | 5.00E+08 | 10GB |
| 8KB | 1 | 11.26622 | 0.00071 | 710.087304 | 1.00E+09 | 5.00E+09 | 10GB |
| | 2 | 22.529044 | 0.000355 | 399.484324 | 1.00E+09 | 5.00E+09 | 10GB |
| | 4 | 48.502913 | 0.000165 | 123.70391 | 1.00E+09 | 5.00E+09 | 10GB |
| | 8 | 2251.98416 | 0.000004 | 2.664317 | 1.00E+09 | 5.00E+09 | 10GB |
| 8MB | 1 | 273.694321 | 0.000029 | 40.190823 | 1.00E+09 | 1.00E+10 | 10GB |
| | 2 | 346.312504 | 0.000023 | 31.763219 | 1.00E+09 | 1.00E+10 | 10GB |
| | 4 | 466.457173 | 0.000017 | 23.582015 | 1.00E+09 | 1.00E+10 | 10GB |
| | 8 | 592.145975 | 0.000014 | 18.5765 | 1.00E+09 | 1.00E+10 | 10GB |
| 80MB | 1 | 477.937483 | 0.000017 | 23.015563 | 1.00E+09 | 1.00E+10 | 10GB |
| | 2 | 462.167589 | 0.000017 | 23.80089 | 1.00E+09 | 1.00E+10 | 10GB |
| | 4 | 629.337257 | 0.000013 | 17.478705 | 1.00E+09 | 1.00E+10 | 10GB |
| | 8 | 658.779646 | 0.000012 | 16.697541 | 1.00E+09 | 1.00E+10 | 10GB |

Table . Performance of read+write
Here 5.00E+8=500000000



Observation and conclusion 1:

- In general, the throughput increases as number of threads increase; the latency decreases as the number of threads decreases. It agrees with the strong scaling performance.

- In general, the throughput increases as the block size increases. My explanation is: as each block size gets smaller, you pay more and more of a penalty for the disk seeks.
- Read operations (both sequential read and random read) speed are obviously faster than read+write operation.
- For small block size (8B), sequential read is faster than random read; while for larger block size, random read is faster than sequential read. My explanation is that rand() function has the advantage to make disk takes less time for seeking, especially as block size gets large.
- The optimal concurrency to get the best performance is thread number 8, because we are using strong scaling, the more number of threads will increase the performance.
- The hardware we are testing is

"device": "sda",

"driver": "mptsas",

"interface": "SCSI",

"vendor": "SEAGATE"

- Based on my performance evaluation, it is HDD.

Experiment 49

- Run IOZONE benchmark
- Since our node RAM size is 4G. As per IOZONE instruction, we should Change the -s to reflect half of your total physical RAM, i.e. 2G.

The screenshot is shown as below:

```

cc@hopess:~/src/current
[cc@hopess current]$ ./iozone -a -s 2g
Iozone: Performance Test of File I/O
Version $Revision: 3.394 $
Compiled for 64 bit mode.
Build: linux

Contributors:William Norcott, Don Capps, Isom Crawford, Kirby Collins
Al Slater, Scott Rhine, Mike Wisner, Ken Goss
Steve Landherr, Brad Smith, Mark Kelly, Dr. Alain CYR,
Randy Dunlap, Mark Montague, Dan Million, Gavin Brebner,
Jean-Marc Zucconi, Jeff Blomberg, Benny Halevy, Dave Boone,
Erik Habbinga, Kris Strecker, Walter Wong, Joshua Root,
Fabrice Bacchella, Zhenghua Xue, Qin Li, Darren Sawyer.
Ben England.

Run began: Sat Oct 7 23:04:05 2017

Auto Mode
File size set to 2097152 KB
Command line used: ./iozone -a -s 2g
Output is in Kbytes/sec
Time Resolution = 0.000001 seconds.
Processor cache size set to 1024 Kbytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.

      KB  reflen  write  rewrite  read  reread  random  random  bkwd  record  stride
      KB  reflen  write  rewrite  read  reread  read  write  read  rewrite  read  fwrite  frewrite  fread  freread
2097152      4  186684  675664  1077166  3576829  2921988  485274  1695937  5996317  3214686  366124  789927  2447393  4079077
2097152      8  125294  1031010  3691765  4993717  4816565  272995  2258555  7281119  4347752  101911  1058324  4686321  5195419
2097152     16  413871  886736  2775541  5847938  4183967  416529  2341198  9558354  4327434  154679  866472  2363755  5390235
2097152     32  317784  1109858  3400619  6419803  5465860  351046  2659224  9608587  6063394  528179  904818  2919206  6388359
2097152     64  192100  1080247  3356899  6273487  4902018  286252  3240146  8293038  4617539  188864  160494  2289407  4926859
2097152    128  241715  106338  2688928  4469262  5451240  138752  2263865  8770694  5789737  166940  576844  2337678  4257321
2097152    256  180224  443726  4098753  5353154  5368545  87666  2797377  7542002  5728044  512775  556446  3642323  5694525
2097152    512  143395  817665  5044421  5751611  5263922  552229  2261253  8975964  4291753  284888  175999  2541765  5408527
2097152   1024  105704  148701  2442555  4632134  5934349  155441  2672089  8669784  5972544  149986  565119  2120546  5723780
2097152   2048  892955  358467  2390707  5966766  5309313  1082634  2221187  8262059  4129105  475411  146109  2673046  6020260
2097152   4096  332708  97288  2351328  5797656  5827721  228217  1938185  8576435  3688862  137554  961691  3248834  6239240
2097152   8192  329997  108407  2523117  5718080  5324205  576495  2061779  7937381  4067036  93781  164560  2484991  6108271
2097152  16384  167804  272699  2410189  4198939  4974752  195364  3641685  4607294  4366483  100946  314747  4073683  4748500

iozone test complete.
[cc@hopess current]$

```

Observation and conclusion 2:

The IOZONE benchmark has a better performance than mine. We can improve our implementation, such as look for better disk access APIs, or increase thread number and block size.

Network Benchmark

- The program finds throughput in Mb per second and latency in ms
- We use socket programming and client-server protocol type to do the implementation.
- The total time difference between start and end time is calculated round trip ping-pong message transferring between client and server
- The actual operation data size is adjusted during experiment to make sure the run time is not that long
- The throughput performance is measured using formula

$$\text{Throughput} = \text{total transferring data size} / \text{total time}$$

- The latency performance is measured using formula

$$\text{Latency} = \text{total time} / \text{total times of round trip} \text{ (i.e. latency is the RTT per professor's instruction)}$$

- We only tested loopback performance due to limited resource on Chameleon

System Configuration

Chameleon Cloud(openstack KVM)

Processor: Intel Xeon E312xx (Sandy Bridge) @ 3.09 GHz
1 processor, 32 cores

OS:CentOS7

RAM: 4GB

Disk: 40GB

Experiment 1~16

- TCP program is run for testing data transferring speed under TCP protocol between client and server. Data size is fixed for each thread to ensure strong scaling. We didn't set data size too large, otherwise the running time would be long. Also it will not affect the result since throughput and latency are unit performance measurement. Num_loops in below table is the number of round trips.



Result shows in below table and graphs.

| thread | throughput (Mb/sec) | latency(ms/RT) | Num_loops | datasize |
|--------|------------------------|----------------|-----------|----------|
| 1 | 576.492249 | 0.222032 | 200 | 25.6MB |
| 2 | 6497.961426 | 0.01998 | 200 | 25.6MB |
| 4 | 4259.041504 | 0.030054 | 200 | 25.6MB |
| 8 | 5694.759277 | 0.02247 | 200 | 25.6MB |

Table TCP loopback performance

| thread | throughput (Mb/sec) | latency(ms/RTT) | Num_loops | datasize |
|--------|------------------------|-----------------|-----------|----------|
| 1 | 2353.336182 | 0.054391 | 200 | 25.6MB |
| 2 | 2122.10791 | 0.060317 | 200 | 25.6MB |
| 4 | 2686.520996 | 0.047645 | 200 | 25.6MB |
| 8 | 2354.763184 | 0.054358 | 200 | 25.6MB |

Table UDP loopback performance





- To see if the transferring data size has effect on the performance, we increase the transferring data size for only UDP to save cloud resource. Since UDP is faster than TCP.

More results are shown below:

| thread | throughput (Mb/sec) | latency(ms/RTT) | Num_loops | datasize |
|--------|------------------------|-----------------|-----------|----------|
| 1 | 5644.057617 | 0.022679 | 62500 | 8GB |
| 2 | 8486.951172 | 0.015082 | 62500 | 8GB |
| 4 | 8468.522461 | 0.015115 | 62500 | 8GB |
| 8 | 2745.610352 | 0.04662 | 62500 | 8GB |
| 1 | 4991.004395 | 0.025646 | 1000 | 128MB |
| 2 | 8383.00293 | 0.015269 | 1000 | 128MB |
| 4 | 8145.578125 | 0.015714 | 1000 | 128MB |
| 8 | 3712.01001 | 0.034483 | 1000 | 128MB |

Table more UDP performance with larger data size





Observation and conclusion 1:

- In general, with small transfer data size, the speed of TCP and UDP does not have big difference; with larger data size, UDP is much faster than TCP (since it takes long to get a result if we using TCP to transfer 8GB data, so we only test UDP for larger data size). My explanation is because UDP's nonexistent acknowledge packet (ACK) that permits a continuous packet stream, instead of TCP that acknowledges a set of packets, calculated by using the TCP window size and round-trip time (RTT).
- In general, both TCP and UDP speed increases as threads increases. Since we use strong scaling. But for larger scale data transfer with UDP, larger thread number does not help transfer faster. My explanation is: That's caused by a lock contention on the UDP receive buffer side. Since both threads are using the same socket descriptor, they spend a disproportionate amount of time fighting for a lock around the UDP receive buffer.
- There is not much difference of transferring speed as data size becomes larger, since throughput and latency are both unit performance measurement.

Experiment 17

- Run iperf benchmark for loopback interface

The screen short is shown as below:

```

Last login: Sun Oct 8 17:22:35 2017
[cc@hope11 ~]$
[cc@hope11 ~]$
[cc@hope11 ~]$ ls
iperf3-3.1.3-1.fc24.x86_64.rpm
[cc@hope11 ~]$ rpm -Uvh iperf###.rpm
error: open of iperf###.rpm failed: No such file or directory
[cc@hope11 ~]$
[cc@hope11 ~]$ rpm -Uvh iperf###.rpmiperf3-3.1.3-1.fc24.x86_64.rpm
error: open of iperf###.rpmiperf3-3.1.3-1.fc24.x86_64.rpm failed: No such file or directory
[cc@hope11 ~]$ rpm -Uvh iperf3-3.1.3-1.fc24.x86_64.rpm
error: can't create transaction lock on /var/lib/rpm/.rpm.lock (Permission denied)
[cc@hope11 ~]$ sudo rpm -Uvh iperf3-3.1.3-1.fc24.x86_64.rpm
Preparing...                               ##### [100%]
Updating / installing...
 1:iperf3-3.1.3-1.fc24                      ##### [100%]
[cc@hope11 ~]$ sudo iperf3 -s -p 80
-----
Server listening on 80
-----
Accepted connection from 129.114.33.165, port 44830
[ 5] local 192.168.0.71 port 80 connected to 129.114.33.165 port 44832
[ ID] Interval           Transfer     Bandwidth
[ 5] 0.00-1.00      sec    109 MBytes   915 Mbits/sec
[ 5] 1.00-2.00      sec    131 MBytes   1.09 Gbits/sec
[ 5] 2.00-3.00      sec    129 MBytes   1.08 Gbits/sec
[ 5] 3.00-4.00      sec    131 MBytes   1.10 Gbits/sec
[ 5] 4.00-5.00      sec    123 MBytes   1.03 Gbits/sec
[ 5] 5.00-6.00      sec    127 MBytes   1.07 Gbits/sec
[ 5] 6.00-7.00      sec    121 MBytes   1.01 Gbits/sec
[ 5] 7.00-8.00      sec    118 MBytes   988 Mbits/sec
[ 5] 8.00-9.00      sec    125 MBytes   1.05 Gbits/sec
[ 5] 9.00-10.00     sec    122 MBytes   1.03 Gbits/sec
[ 5] 10.00-10.04    sec    4.16 MBytes   894 Mbits/sec
-----
[ ID] Interval           Transfer     Bandwidth
[ 5] 0.00-10.04     sec    0.00 Bytes   0.00 bits/sec             sender
[ 5] 0.00-10.04     sec    1.21 GBytes 1.04 Gbits/sec            receiver
-----
Server listening on 80
-----

```

```

[cc@hope11 ~]$ sudo iperf3 -c 129.114.33.165 -i 1 -t 10 -p 80
Connecting to host 129.114.33.165, port 80
[ 4] local 192.168.0.71 port 44832 connected to 129.114.33.165 port 80
[ ID] Interval      Transfer    Bandwidth   Retr  Cwnd
[ 4]  0.00-1.00  sec    117 MBytes  985 Mbits/sec  140   641 KBytes
[ 4]  1.00-2.00  sec    130 MBytes  1.09 Gbits/sec   0   779 KBytes
[ 4]  2.00-3.00  sec    130 MBytes  1.09 Gbits/sec   0   899 KBytes
[ 4]  3.00-4.00  sec    131 MBytes  1.10 Gbits/sec  59   615 KBytes
[ 4]  4.00-5.00  sec    123 MBytes  1.03 Gbits/sec   0   754 KBytes
[ 4]  5.00-6.00  sec    127 MBytes  1.07 Gbits/sec  47   631 KBytes
[ 4]  6.00-7.00  sec    120 MBytes  1.01 Gbits/sec   0   762 KBytes
[ 4]  7.00-8.00  sec    118 MBytes  990 Mbits/sec  14   626 KBytes
[ 4]  8.00-9.00  sec    125 MBytes  1.05 Gbits/sec   0   768 KBytes
[ 4]  9.00-10.00 sec    122 MBytes  1.02 Gbits/sec   0   881 KBytes

- - - - -
[ ID] Interval      Transfer    Bandwidth   Retr
[ 4]  0.00-10.00  sec    1.21 GBytes  1.04 Gbits/sec  260
[ 4]  0.00-10.00  sec    1.21 GBytes  1.04 Gbits/sec

sender
receiver

iperf Done.
[cc@hope11 ~]$

```

Observation and Conclusion 2:

- As we can see, the iperf transfer speed is faster than my implementation. My explanation here is, for latency, we use ping utility measurement to measure the RTT. That means we measure latency for data transferred per round trip; while the iperf measures data transferred per second.
- The theoretical memory performance we got from part 2 is that Stdlib copy is 1.96GB per second, comparing with my highest testing throughput which is 8486.52 Mb/s, we get about $8486.52\text{Mb}/1.96\text{GB}=54.28\%$ performance of the theoretical memory performance. It's reasonable.
- The theoretical loopback is faster than 40Gbps, so the theoretical loopback latency is at least $1/(40\text{Gb}/64\text{KB}/4) = 0.0000256$ seconds = 0.0256ms; Comparing this theoretical latency with my testing latency, they are pretty close. Some of mine are a little faster due to multi threads. Some of mine are a little slower due to loopback network traffic.